

# CS259 - LAB3

The objective of this lab is to introduce you to the basics of extending GPU microarchitecture to accelerate a kernel in hardware

We will be using the Vortex OpenGPU (<https://github.com/vortexgpgpu/vortex>) for this experiment.

You will be using the latest code from branch: **develop**

## Part1 - ISA Extension (1 points)

You will add a new RISC-V custom instruction for computing the integer dot product; **VX\_DOT8**.

VX\_DOT8 calculates the dot product of two small vectors of int8 integers.

Each source register (rs1 and rs2) holds four int8 elements. The operation proceeds as follows:

Extract four int8 values from each source register:

rs1 holds A1, A2, A3, A4

rs2 holds B1, B2, B3, B4

Perform the dot product operation:

Dot Product =  $(A1*B1 + A2*B2 + A3*B3 + A4*B4)$

The result, which could potentially be a 32-bit integer due to overflow, is stored in the destination register (rd).

Use the R-Type RISC-V instruction format.

funct7	rs2	rs1	funct3	rd	opcode
7 bits	5 bit	5 bit	3 bits	5 bit	7 bits

- opcode: Specific opcode reserved for custom instructions.
- rs1, rs2: Source registers containing vectors of four `int8` values each.
- rd: Destination register.
- funct3 and funct7: Used for further specifying the operation (like saturation and accumulation options).

Use custom extension opcode=0x0B with func7=1 and func3=0;

Your new instruction will be implemented as a new operation for the ALU unit.

You will need to modify `vx_intrinsics.h` to define your new VX\_DOT8 instruction

Read the following doc to understand **insn** pseudo instruction format

[https://sourceware.org/binutils/docs/as/RISC\\_002dV\\_002dFormats.html](https://sourceware.org/binutils/docs/as/RISC_002dV_002dFormats.html)

```
// DOT8
inline int vx_dot8(int a, int b) {
    size_t ret;
    asm volatile (".insn r ?, ?, ?, ?, ?, ?" : "=r"(?) : "i"(?), "r"(?),
"r"(?));
    return ret;
}
```

We recommend checking out how VX\_SPLIT and VX\_PRED instructions are decoded in SimX as reference.

## Part2 - Sample matrix multiplication (2 points)

Implement a simple matrix multiplication GPU code that uses your new H/W extension. Here is a basic CPU implementation example:

```
void matrixMultiply(int8_t A[][N], int8_t B[][N], int32_t C[][N], int N) {
    int i, j, k;
    int packedA, packedB;
    for (i = 0; i < N; i++) {
        for (j = 0; j < N; j++) {
            C[i][j] = 0;
            for (k = 0; k < N; k += 4) {
                // Pack 4 int8_t elements from A and B into 32-bit integers
                packedA = *((int*)(A[i] + k));
                packedB = *((int*)(B[k] + j)
                    | *((int*)(B[k+1] + j) << 8)
                    | *((int*)(B[k+2] + j) << 16)
                    | *((int*)(B[k+3] + j) << 24);
                // Accumulate the dot product result into the C matrix
                C[i][j] += vx_dot8(packedA, packedB);
            }
        }
    }
}
```

- Clone sgemm test under tests/regression/sgemm into a new folder tests/regressions/dot8.
- Set PROJECT name to dot8 in tests/regressions/dot8/Makefile
- Update matmul\_cpu in main.cpp to operate on int8\_t matrices.
- Update kernel\_body in tests/regressions/dot8/kernel.cpp to use vx\_dot8

### Part3 - SimX implementation (3 points)

Modify the cycle level simulator to implement the custom ISA extension.

We recommend checking out how VX\_SPLIT and VX\_PRED instructions are decoded in SimX as reference.

Update op\_string() in decode.cpp to print out the new instruction.

Update Emulator::decode() in decode.cpp to decode the new instruction format.

```
switch (func7) {
case 1:
    switch (func3) {
    case 0: // DOT8
        instr->setDestReg(rd, RegType::Integer);
        instr->addSrcReg(rs1, RegType::Integer);
        instr->addSrcReg(rs2, RegType::Integer);
        Break;
```

Update AluType enum in types.h to add DOT8 type

Update Emulator::execute() in execute.cpp to implement the actual VX\_DOT8 emulation.

You will execute the new instruction on the ALU functional unit.

```
switch (func7) {
case 1:
    switch (func3) {
    case 0: { // DOT8
        trace->fu_type = FUType::ALU;
        trace->alu_type = AluType::DOT8;
        trace->used_iregs.set(rsrc0);
        trace->used_iregs.set(rsrc1);
        for (uint32_t t = thread_start; t < num_threads; ++t) {
            if (!warp.tmask.test(t))
                continue;
            // TODO:
        }
        rd_write = true;
    } break;
} break;
}
```

Update AluUnit::tick() in func\_unit.cpp to implement the timing of VX\_DOT8.  
You will assume 2 cycles latency for the dot-product execution.

```
case AluType::DOT8:
```

```
// TODO:
```

```
break;
```

#### **Part4 - RTL implementation (4 points)**

Modify the RTL code to implement the custom ISA extension.

You will need to modify VX\_decode.sv, VX\_alu\_unit.sv primarily.

You will add a new module called VX\_alu\_dot8 and add it VX\_alu\_unit.sv similar to how we added aVX\_alu\_muldiv.

You will assume 2 cycles latency for the operation.

We recommend checking out how MULT instructions are decoded in RTL as reference.

Modify the RTL code to implement the custom ISA extension. We recommend checking out how MULT instructions are decoded and executed in RTL as reference.

Update hw/rtl/VX\_define.vh to define INST\_ALU\_DOT8 as 4'b0001

Update hw/rtl/VX\_config.vh to define LATENCY\_DOT8 as 2

Update dpi\_trace() in hw/rtl/VX\_gpu\_pkg.sv to print the new instruction

Update hw/rtl/core/VX\_decode.sv to decode the new instruction.

Select the ALU functional unit for executing this new instruction.

```
7'h01: begin
```

```
case (func3)
```

```
3'h0: begin // DOT8
```

```
ex_type = // TODO: destination functional unit
```

```
op_type = // TODO: instruction type
```

```
use_rd = // TODO: writing back to rd
```

```
// TODO: set using rd
```

```
// TODO: set using rs1
```

```
// TODO: set using rs2
```

```
end
```

```
default::;
```

```
endcase
```

```
End
```

Create a new VX\_alu\_dot8.sv module that implements DOT8

```
`include "VX_define.vh"

module VX_alu_dot8 #(
    parameter `STRING INSTANCE_ID = "",
    parameter NUM_LANES = 1
) (
    input wire      clk,
    input wire      reset,

    // Inputs
    VX_execute_if.slave execute_if,

    // Outputs
    VX_commit_if.master commit_if
);
    localparam PID_BITS = `CLOG2(`NUM_THREADS / NUM_LANES);
    localparam PID_WIDTH = `UP(PID_BITS);
    localparam TAG_WIDTH = `UUID_WIDTH + `NW_WIDTH + NUM_LANES + `XLEN +
`NR_BITS + 1 + PID_WIDTH + 1 + 1;
    localparam LATENCY_DOT8 = `LATENCY_DOT8;
    localparam PE_RATIO = 2;
    localparam NUM_PES = `UP(NUM_LANES / PE_RATIO);

    `UNUSED_VAR (execute_if.data.op_type)
    `UNUSED_VAR (execute_if.data.op_mod)
    `UNUSED_VAR (execute_if.data.use_PC)
    `UNUSED_VAR (execute_if.data.use_imm)
    `UNUSED_VAR (execute_if.data.tid)
    `UNUSED_VAR (execute_if.data.rs3_data)

    wire [NUM_LANES-1:0][2*`XLEN-1:0] data_in;

    for (genvar i = 0; i < NUM_LANES; ++i) begin
        assign data_in[i][0+: `XLEN] = execute_if.data.rs1_data[i];
        assign data_in[i][`XLEN+: `XLEN] = execute_if.data.rs2_data[i];
    end

    wire pe_enable;
    wire [NUM_PES-1:0][2*`XLEN-1:0] pe_data_in;
    wire [NUM_PES-1:0][`XLEN-1:0] pe_data_out;

    // PEs time-multiplexing
```

```

VX_pe_serializer #(
    .NUM_LANES (NUM_LANES),
    .NUM_PES   (NUM_PES),
    .LATENCY   (LATENCY_DOT8),
    .DATA_IN_WIDTH (2*`XLEN),
    .DATA_OUT_WIDTH (`XLEN),
    .TAG_WIDTH  (TAG_WIDTH),
    .PE_REG     (1)
) pe_serializer (
    .clk      (clk),
    .reset     (reset),
    .valid_in  (execute_if.valid),
    .data_in   (data_in),
    .tag_in    ({
        execute_if.data.uuid,
        execute_if.data.wid,
        execute_if.data.tmask,
        execute_if.data.PC,
        execute_if.data.rd,
        execute_if.data.wb,
        execute_if.data.pid,
        execute_if.data.sop,
        execute_if.data.eop
    }),
    .ready_in  (execute_if.ready),
    .pe_enable (pe_enable),
    .pe_data_in (pe_data_in),
    .pe_data_out(pe_data_out),
    .valid_out  (commit_if.valid),
    .data_out   (commit_if.data.data),
    .tag_out    ({
        commit_if.data.uuid,
        commit_if.data.wid,
        commit_if.data.tmask,
        commit_if.data.PC,
        commit_if.data.rd,
        commit_if.data.wb,
        commit_if.data.pid,
        commit_if.data.sop,
        commit_if.data.eop
    }),
    .ready_out  (commit_if.ready)
);

```

```

// PEs instancing
for (genvar i = 0; i < NUM_PES; ++i) begin
    wire [31:0] a = pe_data_in[i][0 +: 32];
    wire [31:0] b = pe_data_in[i][32 +: 32];
    // TODO:
    wire [31:0] result;
    `BUFFER_EX(result, c, pe_enable, LATENCY_DOT8);
    assign pe_data_out[i] = result;
end

```

**Endmodule**

Update hw/rtl/core/VX\_alu\_unit.sv to add your new VX\_alu\_dot8 instance as a 3rd sub-unit after VX\_alu\_muldiv.

### **Report:**

You will compare your new accelerated sgemm program with the existing sgemm kernel under the regression codebase. You will use  $N=128$  and (warps=4, threads=4) and (Warps=16, threads=16) for 1 and 4 cores on the SimX and RTLSIM driver. Plot the IPC,

### **Bonus (2 points):**

You will get a bonus if you write an optimized kernel that leverages the GPU local memory to tile the computation on 16x16 elements (256 bytes tiles).

### **What to submit?**

1- zip file containing report.pdf and changes.diff.

The report should be about 2 pages showing the generated plots.